

BIRLA INSTITUTE OF TECHNOLOGY & SCIENCE, PILANI (RAJASTHAN)

CS F422 – Parallel Computing

Lab#4

Note: Please use programs under *Code_lab4* directory supplied with this sheet. Do not copy from this sheet.

The lab has the following objectives:

Giving practice programs for CUDA.

Query device features

```
1. #include "./book.h"
2.
3. int main(void)
4. {
5.     cudaDeviceProp prop;
6.     int count;
7.     HANDLE_ERROR(cudaGetDeviceCount(&count));
8.     for (int i = 0; i < count; i++)
9.     {
10.         HANDLE_ERROR(cudaGetDeviceProperties(&prop, i));
11.         printf(" --- General Information for device %d ---\n", i);
12.         printf("Name: %s\n", prop.name);
13.         printf("Compute capability: %d.%d\n", prop.major,
prop.minor);
14.         printf("Clock rate: %d\n", prop.clockRate);
15.         printf("Device copy overlap: ");
16.         if (prop.deviceOverlap)
17.             printf("Enabled\n");
18.         else
19.             printf("Disabled\n");
20.         printf("Kernel execution timeout : ");
21.         if (prop.kernelExecTimeoutEnabled)
22.             printf("Enabled\n");
23.         else
24.             printf("Disabled\n");
25.         printf(" --- Memory Information for device %d ---\n", i);
26.         printf("Total global mem: %ld\n", prop.totalGlobalMem);
27.         printf("Total constant Mem: %ld\n", prop.totalConstMem);
28.         printf("Max mem pitch: %ld\n", prop.memPitch);
29.         printf("Texture Alignment: %ld\n", prop.textureAlignment);
30.         printf(" --- MP Information for device %d ---\n", i);
31.         printf("Multiprocessor count: %d\n",
prop.multiProcessorCount);
32.         printf("Shared mem per mp: %ld\n", prop.sharedMemPerBlock);
33.         printf("Registers per mp: %d\n", prop.regsPerBlock);
34.         printf("Threads in warp: %d\n", prop.warpSize);
35.
```

```

36.         printf("Max threads per block: %d\n",
37.               prop.maxThreadsPerBlock);
38.         printf("Max thread dimensions: (%d, %d, %d)\n",
39.               prop.maxThreadsDim[0], prop.maxThreadsDim[1],
40.               prop.maxThreadsDim[2]);
41.         printf("Max grid dimensions: (%d, %d, %d)\n",
42.               prop.maxGridSize[0], prop.maxGridSize[1],
43.               prop.maxGridSize[2]);
44.         printf("\n");
45.     }
46. }

```

Q?

1. This is a program to query the device features in your GPU. See if you can understand most of what they mean, and if they match your device configuration.

Using parameters to kernel

```

1. #include <iostream>
2. #include "book.h"
3.
4. __global__ void add(int a, int b, int *c)
5. {
6.     *c = a + b;
7. }
8.
9. int main(void)
10. {
11.     int c;
12.     int *dev_c;
13.     HANDLE_ERROR(cudaMalloc((void **)&dev_c, sizeof(int)));
14.     add<<<1, 1>>>(2, 7, dev_c);
15.     HANDLE_ERROR(cudaMemcpy(&c,
16.                             dev_c,
17.                             sizeof(int),
18.                             cudaMemcpyDeviceToHost));
19.     printf("2 + 7 = %d\n", c);
20.     cudaFree(dev_c);
21.     return 0;
22. }

```

Q?

1. This is a program to teach passing of parameters to the kernel.

GPU Vector sum using parallel blocks

```
1. #include "../book.h"
2. #define N 10
3.
4. int main(void)
5. {
6.     int a[N], b[N], c[N];
7.     int *dev_a, *dev_b, *dev_c;
8.     // allocate the memory on the GPU
9.     HANDLE_ERROR(cudaMalloc((void **)&dev_a, N * sizeof(int)));
10.    HANDLE_ERROR(cudaMalloc((void **)&dev_b, N * sizeof(int)));
11.    HANDLE_ERROR(cudaMalloc((void **)&dev_c, N * sizeof(int)));
12.    // fill the arrays 'a' and 'b' on the CPU
13.    for (int i = 0; i < N; i++)
14.    {
15.        a[i] = -i;
16.        b[i] = i * i;
17.    }
18.
19.    // copy the arrays 'a' and 'b' to the GPU
20.    HANDLE_ERROR(cudaMemcpy(dev_a, a, N * sizeof(int),
        cudaMemcpyHostToDevice));
21.    HANDLE_ERROR(cudaMemcpy(dev_b, b, N * sizeof(int),
        cudaMemcpyHostToDevice));
22.
23.    add<<<N,1>>>(dev_a, dev_b, dev_c);
24.
25.    // copy array 'c' back from the GPU to the CPU
26.    HANDLE_ERROR(cudaMemcpy(c, dev_c, N * sizeof(int),
        cudaMemcpyDeviceToHost));
27.
28.    // display the results
29.    for(int i=0; i<N; ++i){
30.        printf("%d + %d = %d\n", a[i], b[i], c[i]);
31.    }
32.
33.    // free the memory allocated on the GPU
34.    cudaFree(dev_a);
35.    cudaFree(dev_b);
36.    cudaFree(dev_c);
37.
38.    return 0;
39. }
```

Q?

1. This is a program which uses GPU to compute the vector sum using a parallel block implementation.

GPU Vector sum using threads

```

1. #include "../common/book.h"
2. #define N 10
3. __global__ void add(int *a, int *b, int *c)
4. {
5.     int tid = threadIdx.x;
6.     if (tid < N)
7.         c[tid] = a[tid] + b[tid];
8. }
9. int main(void)
10. {
11.     int a[N], b[N], c[N];
12.     int *dev_a, *dev_b, *dev_c;
13.     // allocate the memory on the GPU
14.     HANDLE_ERROR(cudaMalloc((void **)&dev_a, N * sizeof(int)));
15.     HANDLE_ERROR(cudaMalloc((void **)&dev_b, N * sizeof(int)));
16.     HANDLE_ERROR(cudaMalloc((void **)&dev_c, N * sizeof(int)));
17.     // fill the arrays "a" and "b" on the CPU
18.     for (int i = 0; i < N; i++)
19.     {
20.         a[i] = i;
21.         b[i] = i * i;
22.     }
23.     // copy the arrays "a" and "b" to the GPU
24.     HANDLE_ERROR(cudaMemcpy(dev_a,
25.                             a,
26.                             N * sizeof(int),
27.                             cudaMemcpyHostToDevice));
28.     HANDLE_ERROR(cudaMemcpy(dev_b,
29.                             b,
30.                             N * sizeof(int),
31.                             cudaMemcpyHostToDevice));
32.     add<<<1, N>>>>(dev_a, dev_b, dev_c);
33.     // copy the array "c" back from the GPU to the CPU
34.     HANDLE_ERROR(cudaMemcpy(c,
35.                             dev_c,
36.                             N * sizeof(int),
37.                             cudaMemcpyDeviceToHost));
38.     // display the results
39.     for (int i = 0; i < N; i++)
40.     {
41.         printf( "% d + % d = % d\n", a[i], b[i], c[i]);
42.     }
43.     // free the memory allocated on the GPU
44.     cudaFree(dev_a);
45.     cudaFree(dev_b);
46.     cudaFree(dev_c);
47.     return 0;
48. }

```

Q?

1. This is a program which uses parallel threads instead of parallel block to compute vector sum using GPU. Can you understand the difference b/w the two programs?

End of lab4